



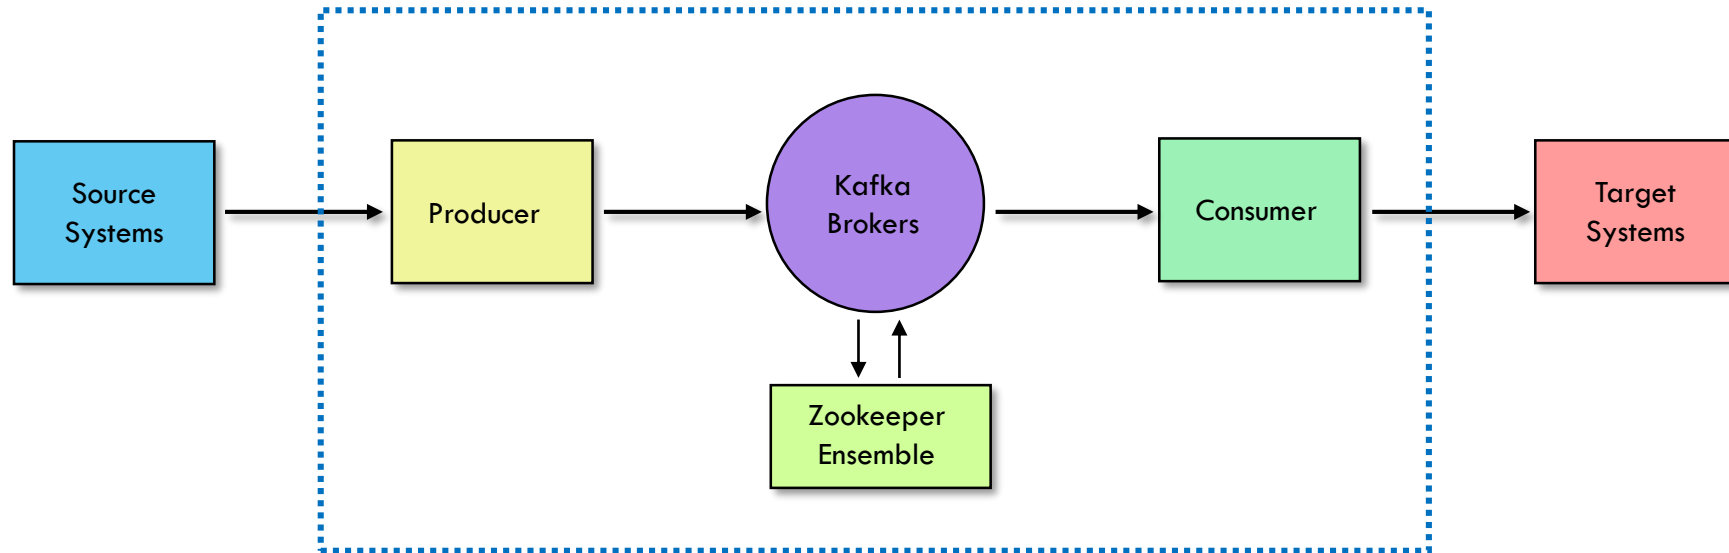
# ak系列 - ak03

## (Java Producer & Consumer - 進階)



# Kafka Ecosystem:

## Kafka Core



# Download Sample Java Project Code

- Get Java Demo source code

 緯創IT先進技術實驗室 (witlab)  主頁

Home / Courses

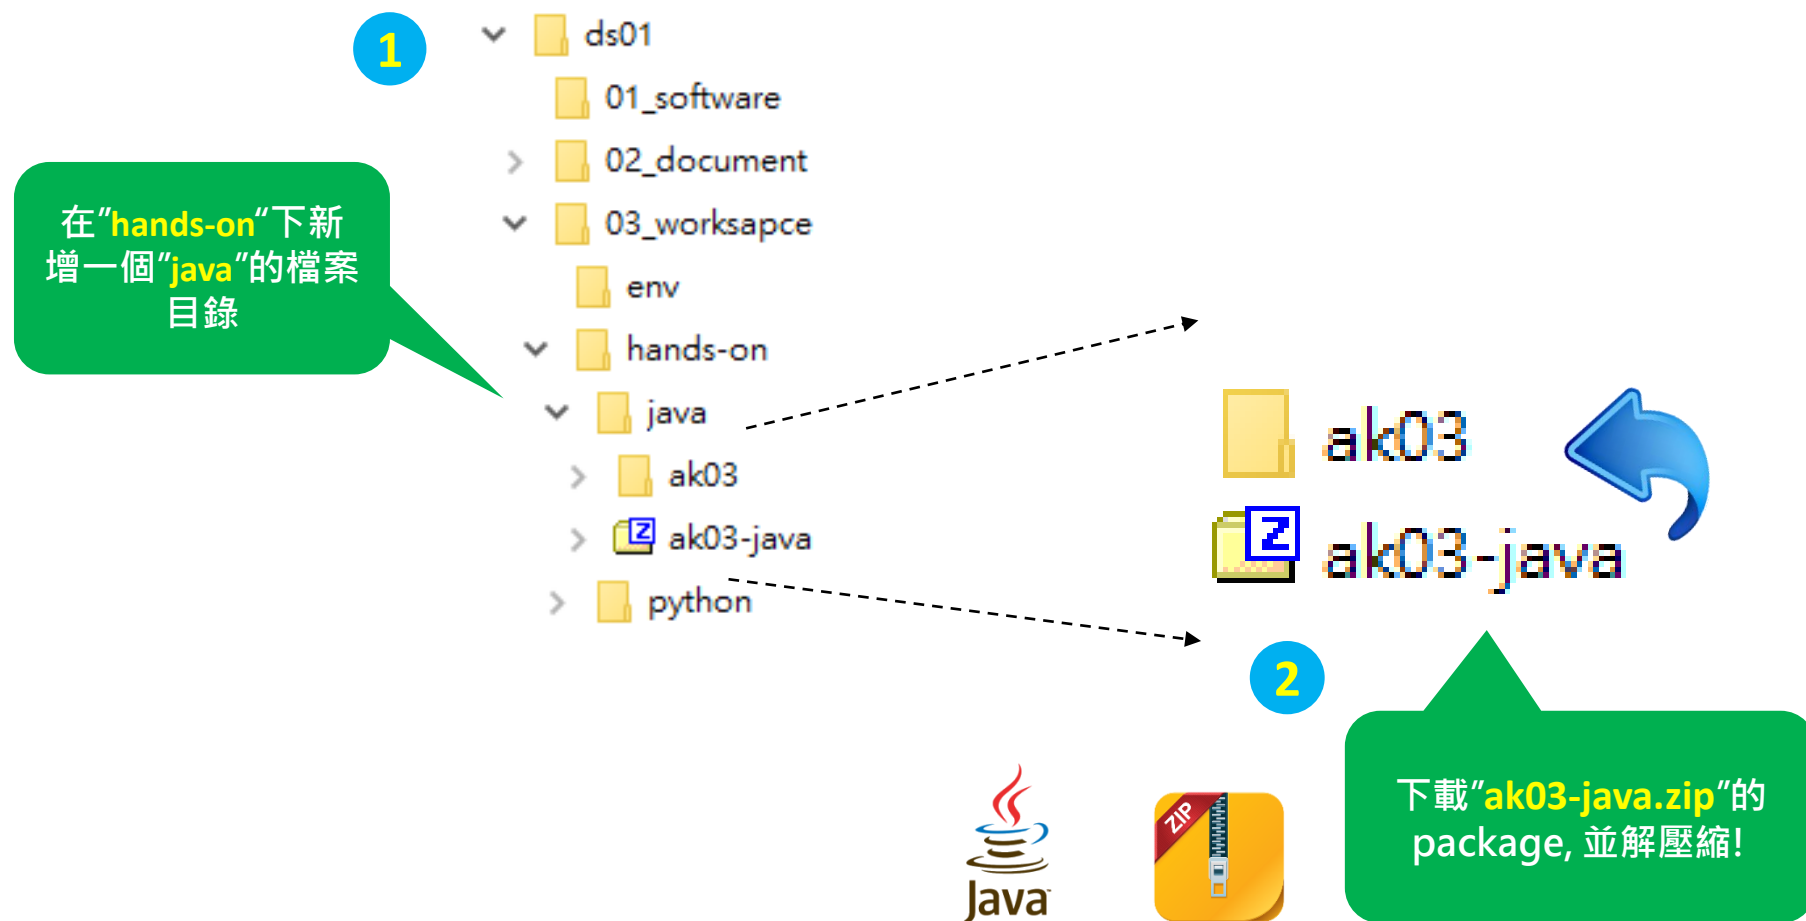
Workshop#  
ds01

課程表

課程單元	課程內容	時間	前提	學員	連結
ak03: Apache Kafka - 訊息發佈與訂閱進階	<ul style="list-style-type: none"><li>• Producer同步與非同步的訊息發佈</li><li>• Producer訊息發佈時的重要參數</li><li>• Producer發佈JSON訊息(序列化概念)</li><li>• Consumer與ConsumerGroup實驗</li><li>• Consumer與Rebalance機制</li><li>• Consumer與Offset的關係</li><li>• Consumer與多種Commit Offset的手法</li><li>• 觀念與實作測驗</li></ul>	3 小時	ak02		<ul style="list-style-type: none"><li>•  Java<ul style="list-style-type: none"><li>◦  簡報 (Java版本)</li><li>◦  簡報 (Java範例) </li><li>◦  作業 (Java範例)</li></ul></li><li>•  Python<ul style="list-style-type: none"><li>◦  簡報 (Python版本)</li><li>◦  簡報 (Python範例)</li><li>◦  作業 (Python範例)</li></ul></li><li>•  簡報 (Scripts)</li><li>•  作業</li><li>•  作業繳交進度</li></ul>



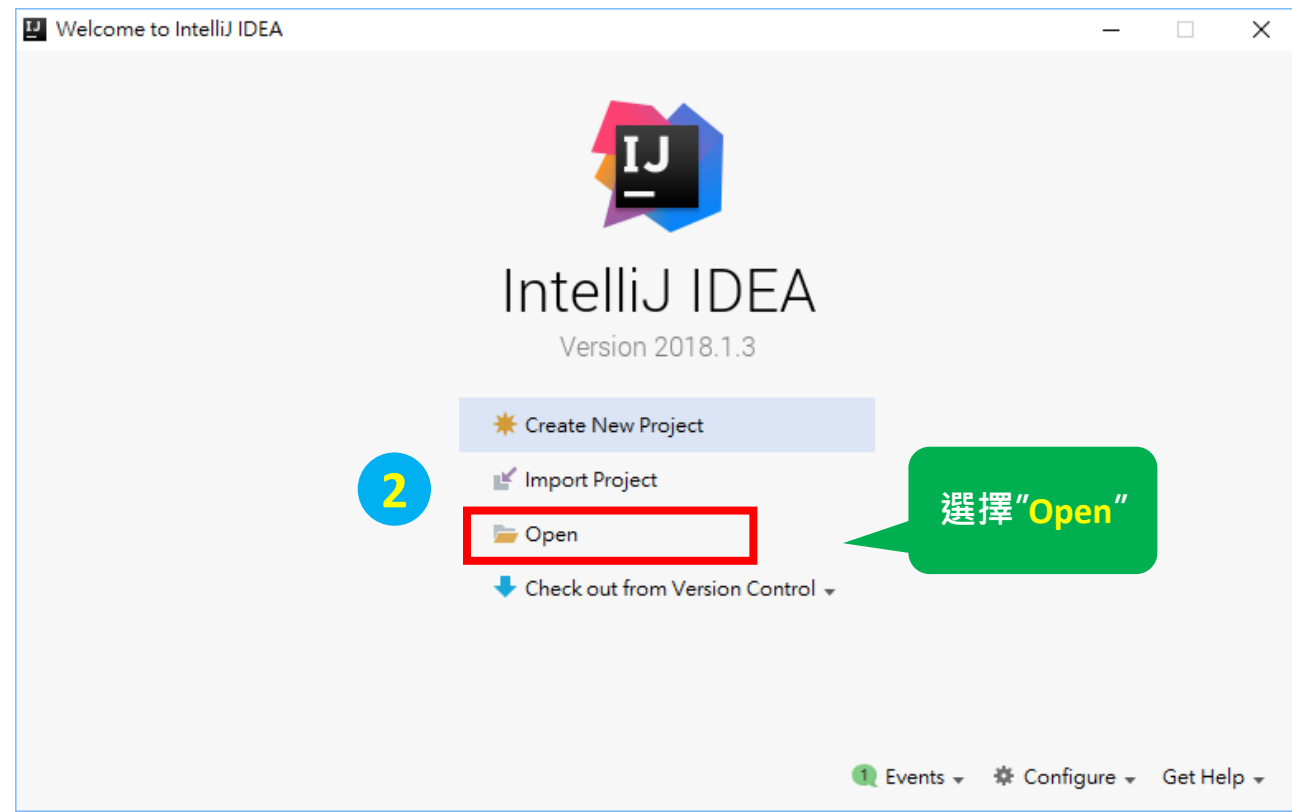
# Decompress Demo Java



# Open Demo Java Project using IntelliJ

1

啟動IntelliJ IDEA




選擇"Open"

# Open Demo Java Project using IntelliJ

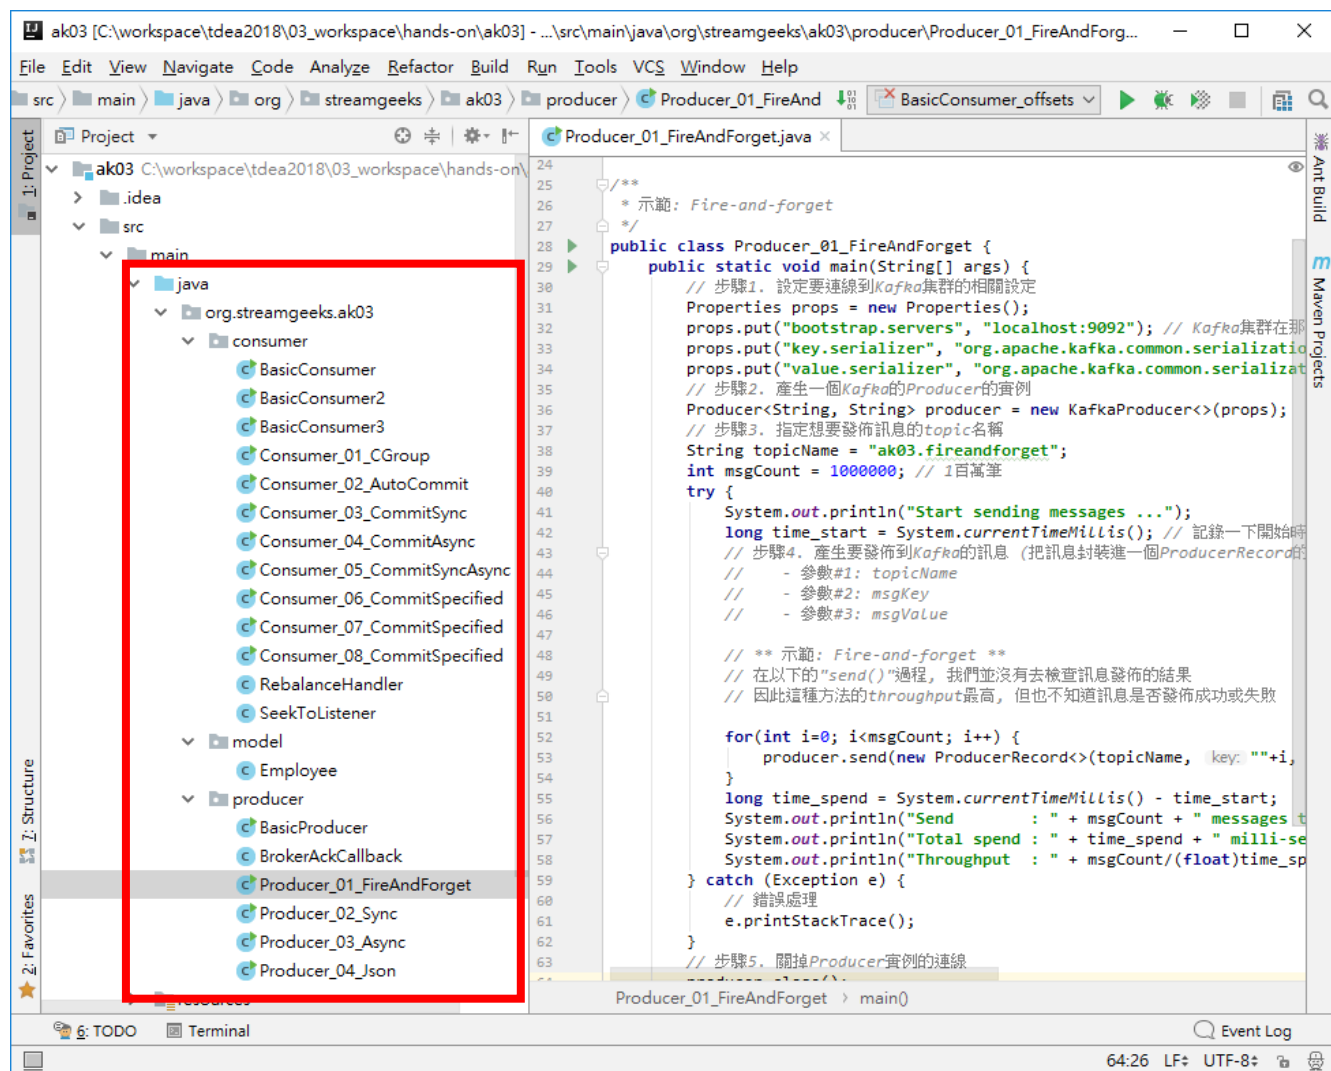
1

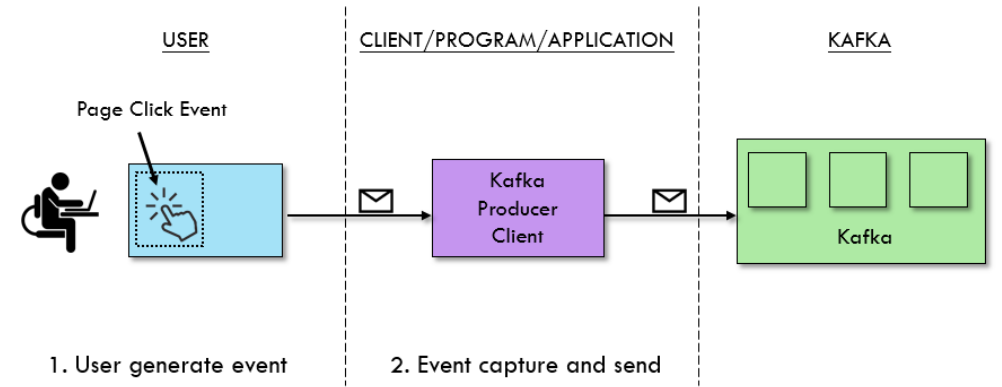
切換到解壓縮的  
目錄



# Open Demo Java Project using IntelliJ

1

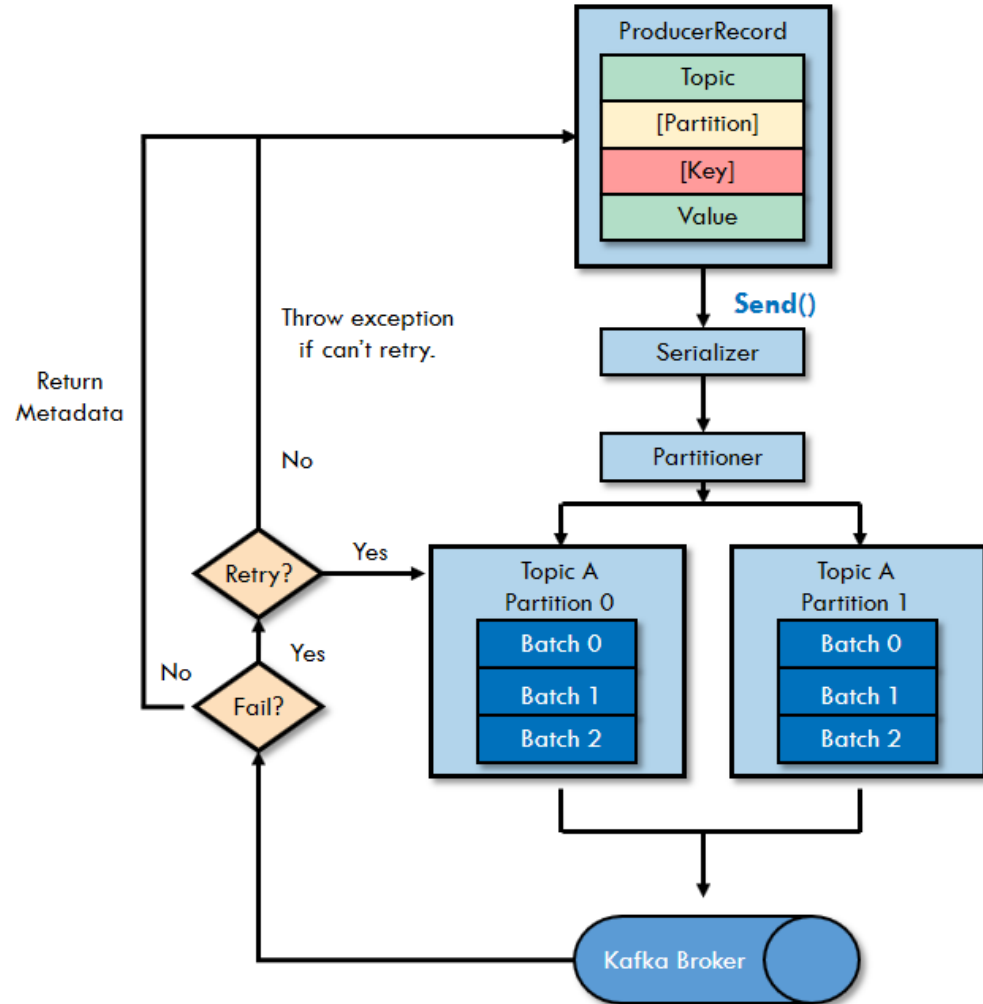
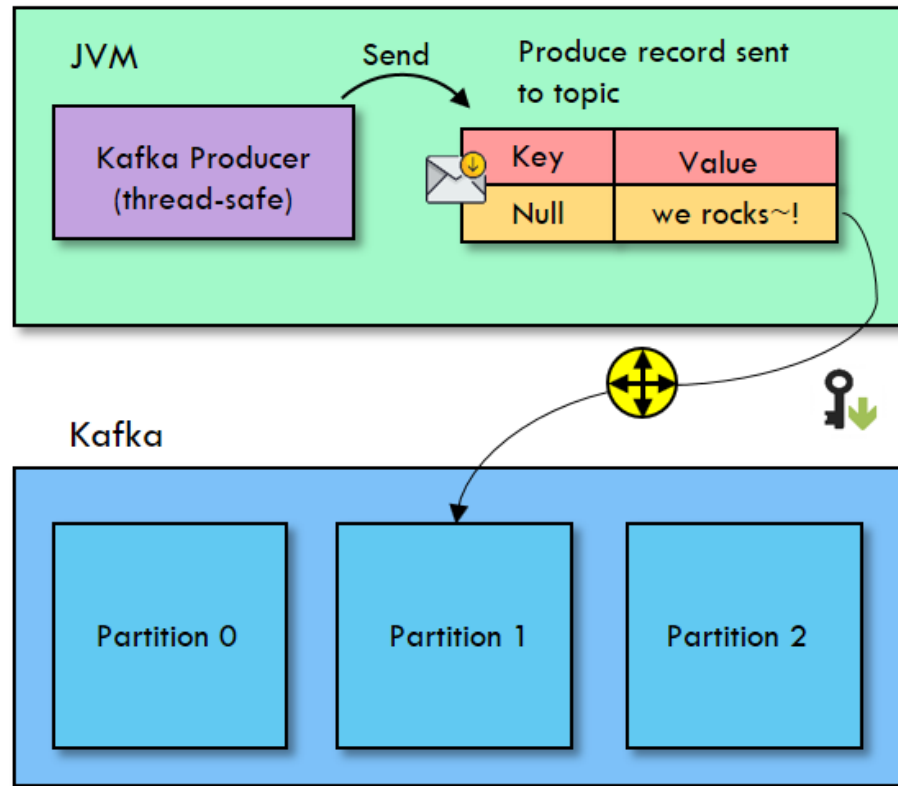




# Kafka Producers: Writing Messages to Kafka



# Producer Overview



# Three primary methods of sending messages

- **Fire-and-forget**

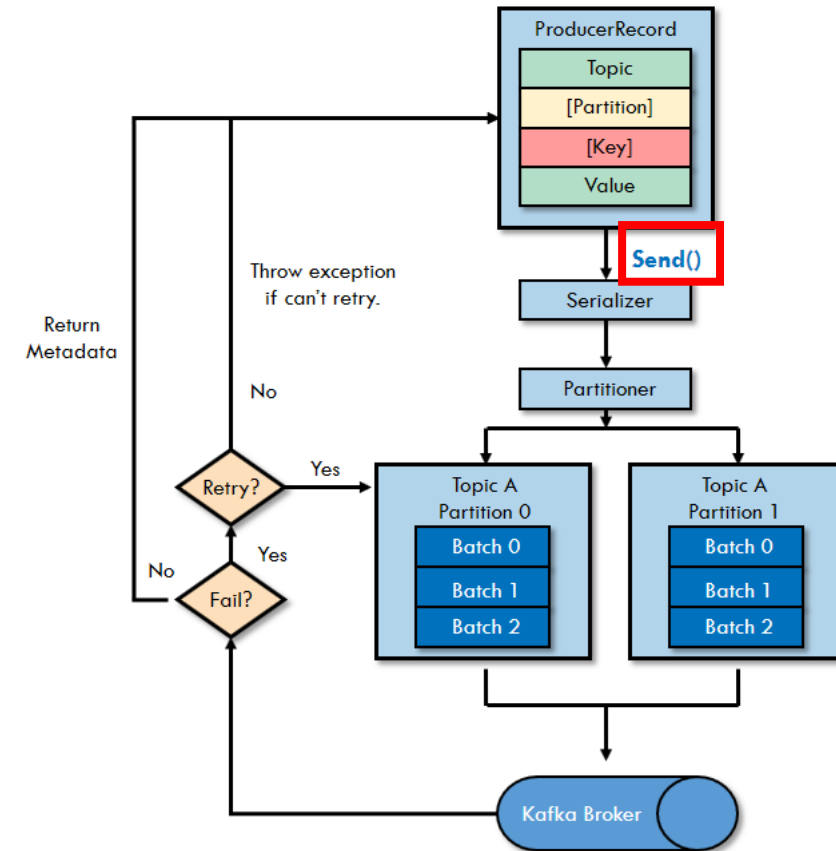
- We send a message to the server and don't really care if it arrives successfully or not.

- **Synchronous send**

- We send a message, and **wait** to see if the **send()** was successful or not.

- **Asynchronous send**

- We call the **send()** method with a **callback function**, which gets triggered when it receives a response from the Kafka broker.

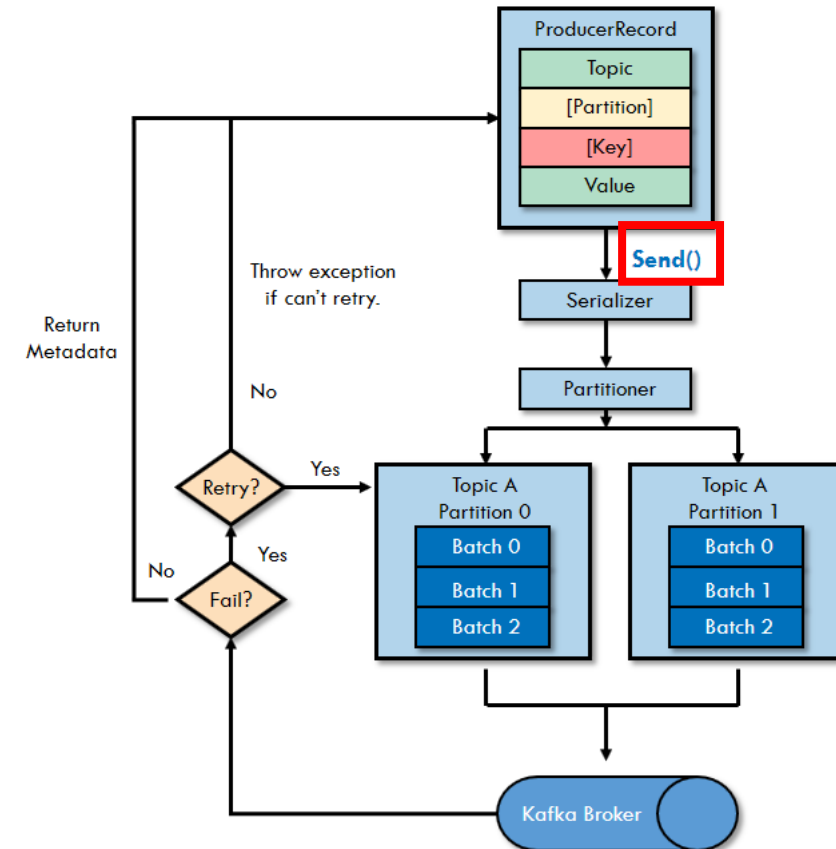


# Three primary methods of sending messages

## Fire-and-forget

- Producer\_01\_FireAndForget

```
// ** 示範: Fire-and-forget **  
// 在以下的"send()"過程, 我們並沒有去檢查訊息發佈的結果  
// 因此這種方法的throughput最高, 但也知道訊息是否發佈成功或失敗  
  
for(int i=0; i<msgCount; i++) {  
    producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i));  
}
```

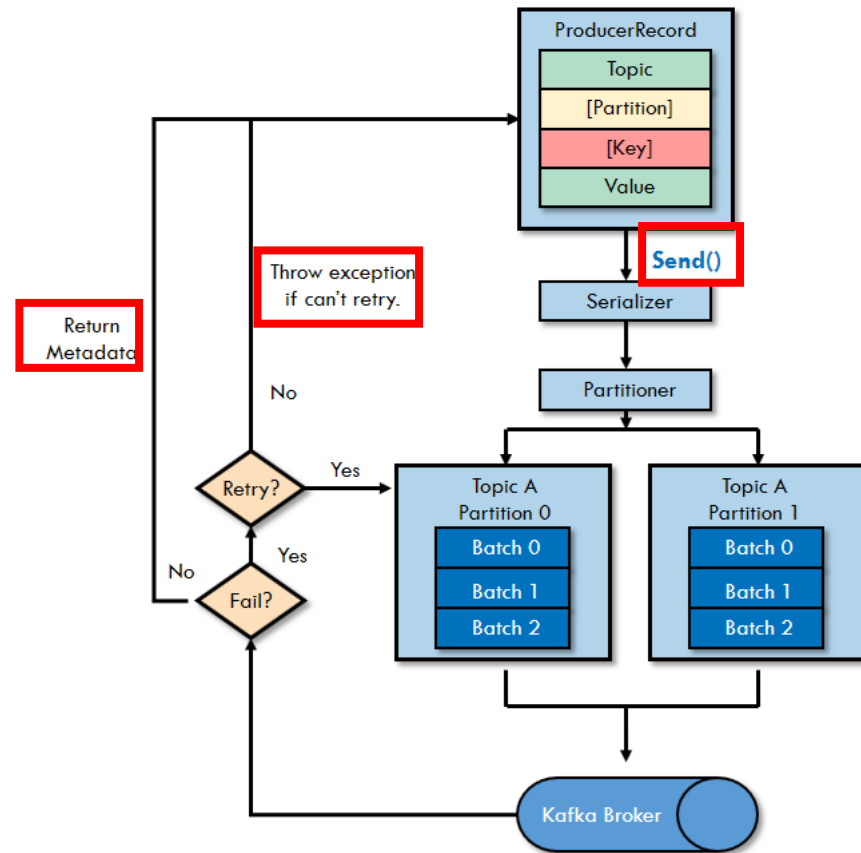


# Three primary methods of sending messages

## Synchronous send

- Producer\_02\_Sync

```
// ** 示範: Synchronous Send **  
// 在以下的"send().get()"的同步發佈訊息的過程，透過同步我們可以取得由Broker回覆訊息發佈的ack結果  
// 由於這種方法是一筆一筆的送與等待Broker的回應。因此這種方法的throughput最差（通常不會用Kafka來做這種事）  
  
for(int i=0; i<msgCount; i++) {  
    // Broker回覆訊息包含了訊息落在topic的那個partition及offset  
    RecordMetadata recordMetadata = producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i)).get();  
    if((i+1)%10000==0) {  
        // 為了不讓打印訊息拖慢速度，我們每1萬打印一筆recordMetadata來看  
        System.out.println((i+1) + " messages sent!");  
        System.out.println("Topic:Partition:Offset: [" + recordMetadata.topic() + "]:["  
            + recordMetadata.partition() + "]:["  
            + recordMetadata.offset() + "]);  
    }  
}
```

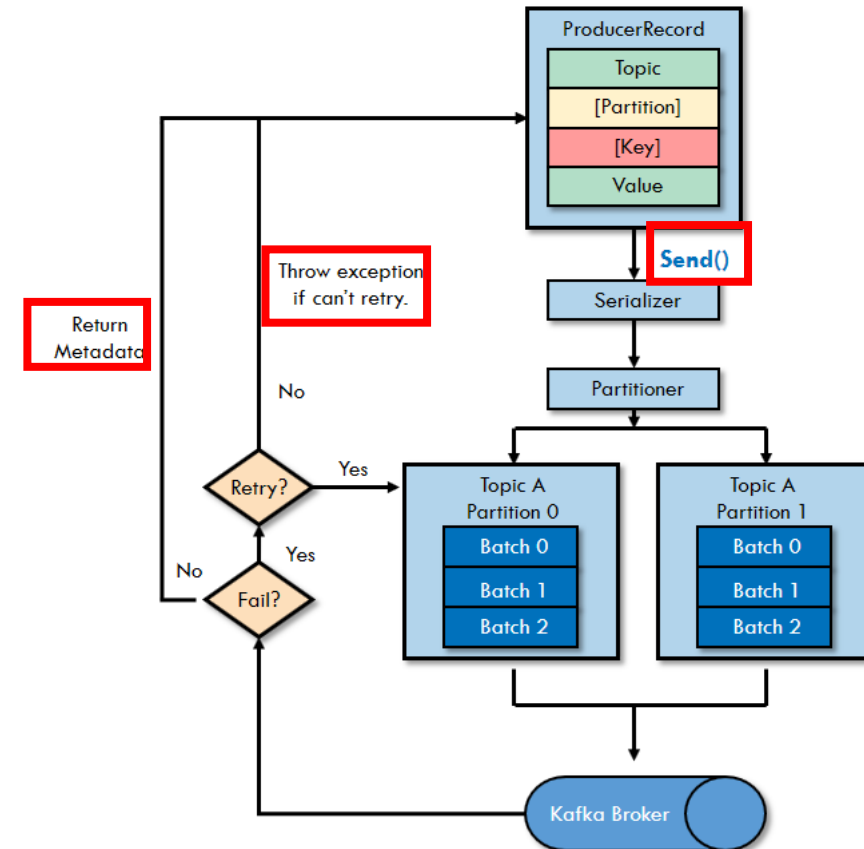


# Three primary methods of sending messages

## Asynchronous send

- Producer\_03\_Async
  - BrokerAckCallback

```
// ** 示範: Asynchronous Send **  
// 透過一個我們自己定義的Callback函式我們可以非同步地取得由Broker回覆訊息發佈的ack結果  
// 這種方法可以取得Broker回覆訊息發佈的ack結果，同時又可以取得好的throughput (建議的作法)  
  
AtomicInteger atomicInteger = new AtomicInteger();  
CountDownLatch countDownLatch = new CountDownLatch(msgCount);  
  
for(int i=0; i<msgCount; i++) {  
    // 回呼函式會在Broker送回ack的時候自動去呼叫  
    producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_" + i),  
        new BrokerAckCallback(atomicInteger, countDownLatch));  
}
```




# Three primary methods of sending messages

## Asynchronous send

- Producer\_03\_Async
  - BrokerAckCallback

```
public class BrokerAckCallback implements Callback {  
    // 使用AtomicInteger來做為非同步的counter, 可以避免blocking及contention  
    AtomicInteger atomicInteger;  
    // 使用CountDownLatch來讓主線程知道所有的訊息都已經處理完了  
    CountDownLatch countDownLatch;  
  
    public BrokerAckCallback(AtomicInteger atomicInteger, CountDownLatch countDownLatch) {  
        this.atomicInteger = atomicInteger;  
        this.countDownLatch = countDownLatch;  
    }  
  
    @Override  
    public void onCompletion(RecordMetadata recordMetadata, Exception exception) {  
        // 增加1再取出  
        int counter = atomicInteger.incrementAndGet();  
        if(counter%100000==0) {  
            // 為了不讓打印訊息拖慢速度, 我們每10萬打印一筆recordMetadata來看  
            System.out.println(counter + " messages sent!");  
            System.out.println("Topic:Partition:Offset: [" + recordMetadata.topic() + "]:["  
                + recordMetadata.partition() + "]:["  
                + recordMetadata.offset() + "]);"  
        }  
        this.countDownLatch.countDown(); // 倒數  
    }  
}
```



# Configuring Producers - Serializers

Kafka lets us publish and subscribe to streams of records and the records can be of any type (JSON, String, POJO, etc.)

- **key.serializer & value.serializer**

- **String – StringSerializer**

- JSON, XML, Delimited Text, etc..

- Integer - IntegerSerializer

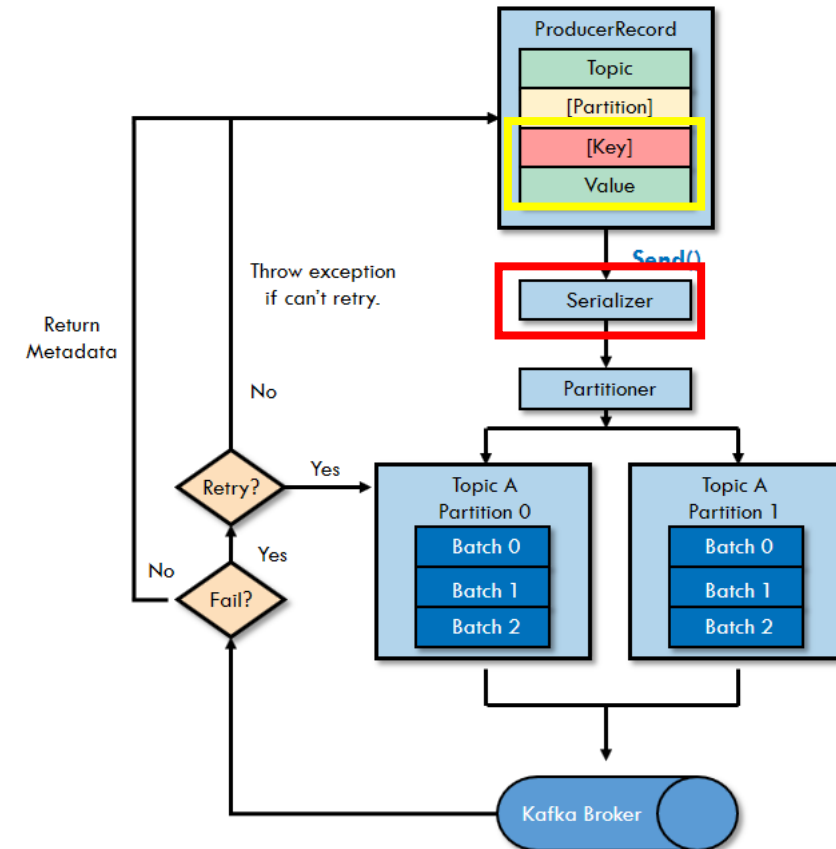
- Long - LongSerializer

- Double - DoubleSerializer

- ByteBuffer - ByteBufferSerializer

- **byte[] – ByteArraySerializer**

- Avro, ProtoBuf, etc..



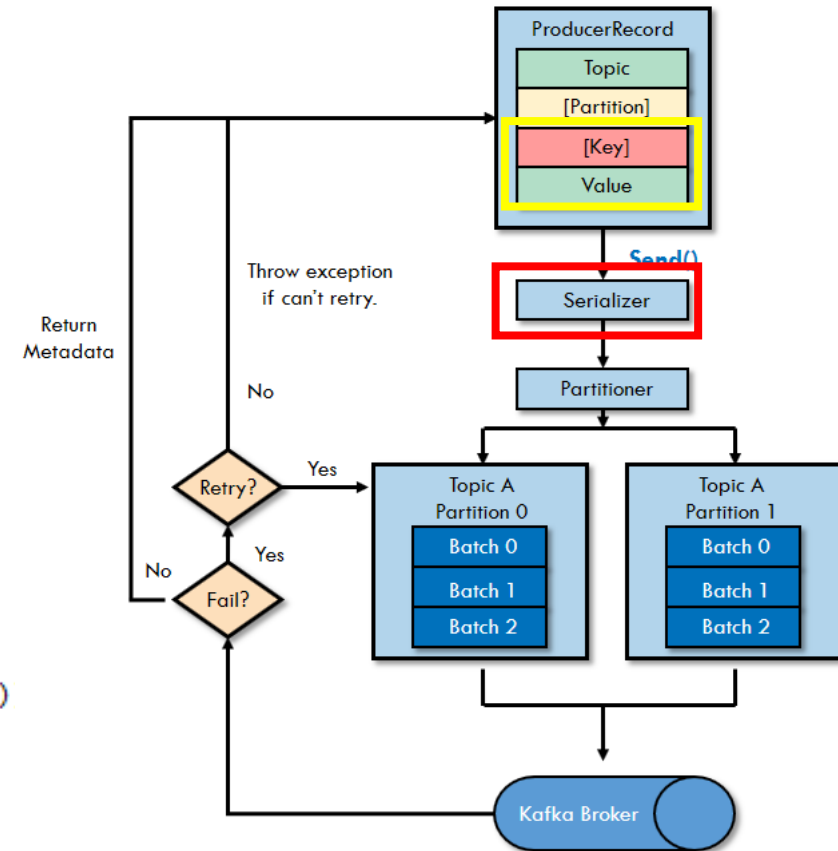
# Configuring Producers - Serializers

- Producer\_04\_Json
  - **key.serializer & value.serializer**
    - String – StringSerializer (JSON)

```
props.put("bootstrap.servers", "localhost:9092"); // Kafka集群在那裡?  
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer"); // 指定msgKey的序列化器  
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer"); // 指定msgValue的序列化器
```

```
// 一個用來將DTO物件轉換成(JSON String)的物件 <-- 重要  
ObjectMapper om = new ObjectMapper();
```

```
for(int i=0; i<msgCount; i++) {  
    // 讓我們產生假的Employee資料  
    Employee employee = new Employee( id: "empid_"+i, firstName: "fn_"+i, lastName: "ln_"+i,  
    |      deptid: "deptid_"+i%10, new Date(), wage: 100000*random.nextFloat(), random.nextBoolean()  
  
    // 把Employee轉換成JSON字串(String)  
    String employeeJSON = om.writeValueAsString(employee);  
  
    // 送出訊息  
    producer.send(new ProducerRecord<>(topicName, employee.getId(), employeeJSON),  
        new BrokerAckCallback(atomicInteger, countDownLatch)); // 回呼函式會在Broker送回ack的時候自動去呼叫
```

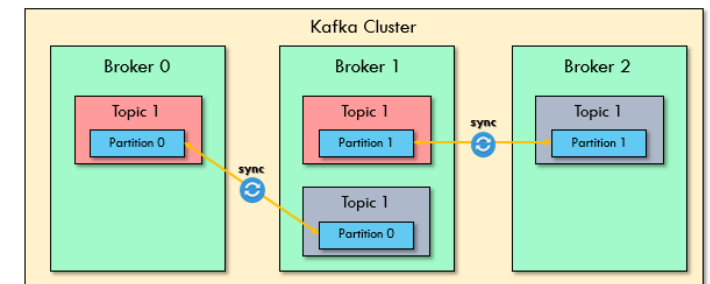
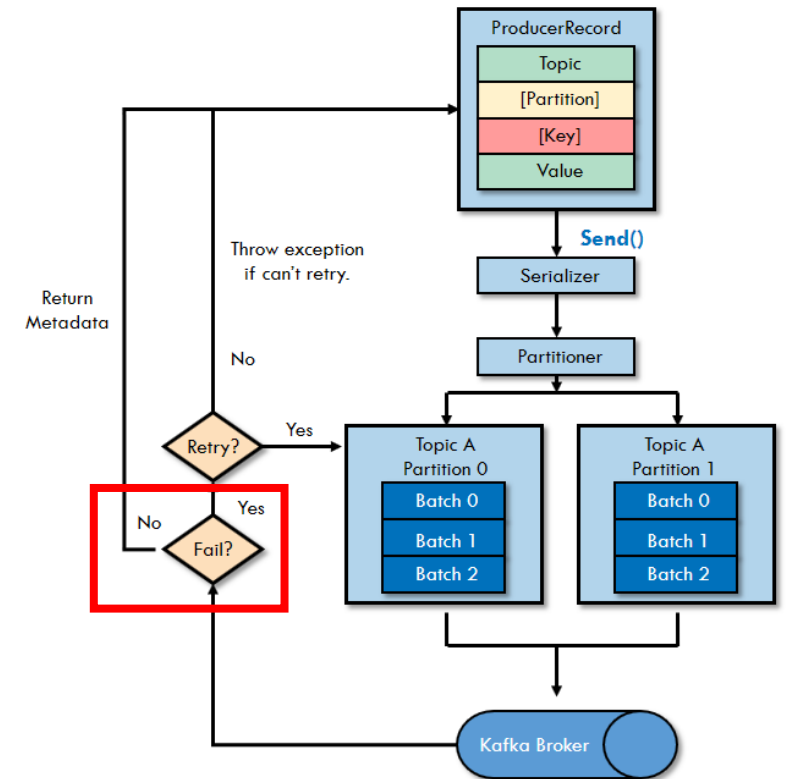




# Configuring Producers

- **acks=0**, the producer will not wait for a reply from the broker before assuming the message was sent successfully.
- **acks=1**, the producer will receive a success response from the broker the moment the leader replica received the message.
- **acks=all**, the producer will receive a success response from the broker once all in-sync replicas received the message.

```
props.put("acks", "all");
```

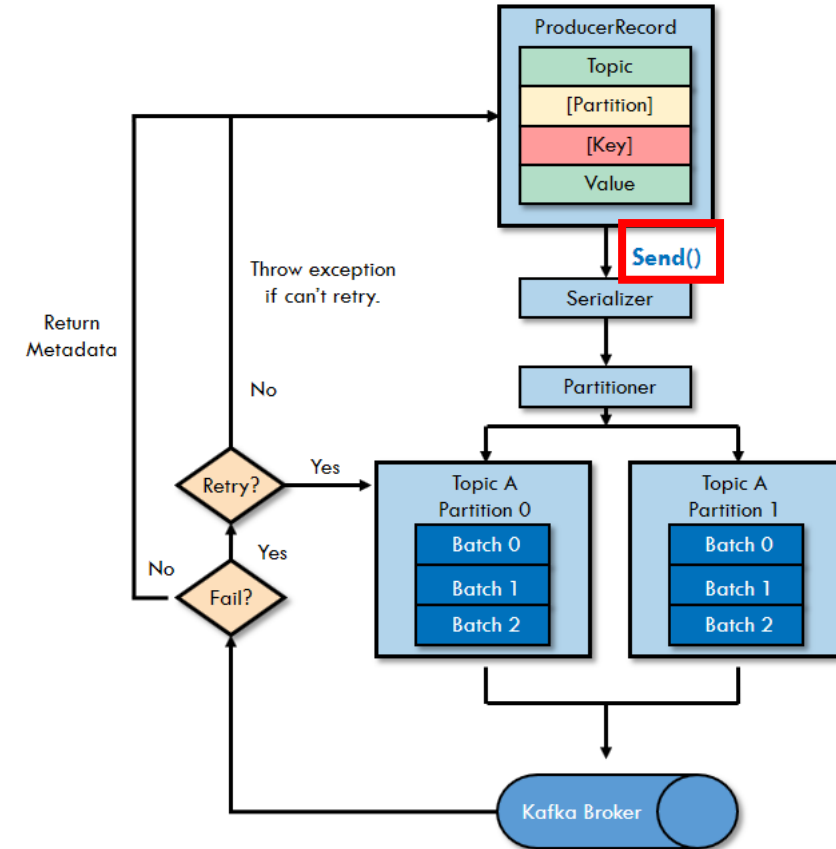


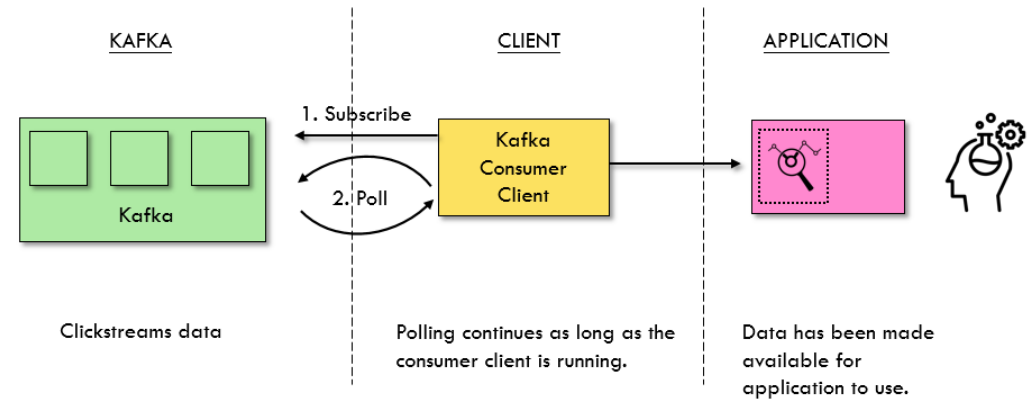
# Configuring Producers

- **enable.idempotence**

- When set to true, the producer will ensure that messages are successfully produced exactly once and in the original produce order.

```
props.put("enable.idempotence", "true");
```

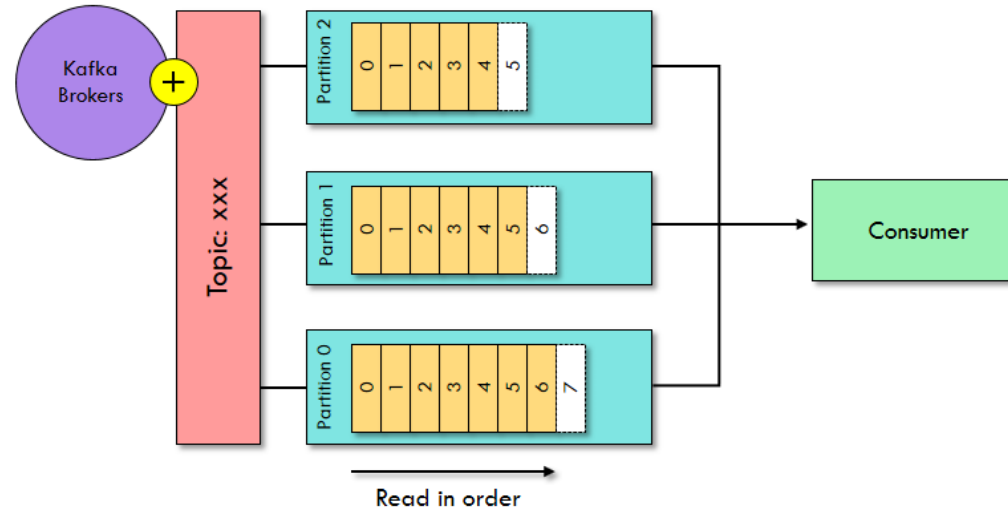




# Kafka Consumers: Reading Messages from Kafka


# Consumers

- **Consumers** read data from a topic
- They only have to specify the topic name and one broker to connect to, and Kafka will automatically take care of pulling the data from the right brokers
- Data is read in order **for each partitions**



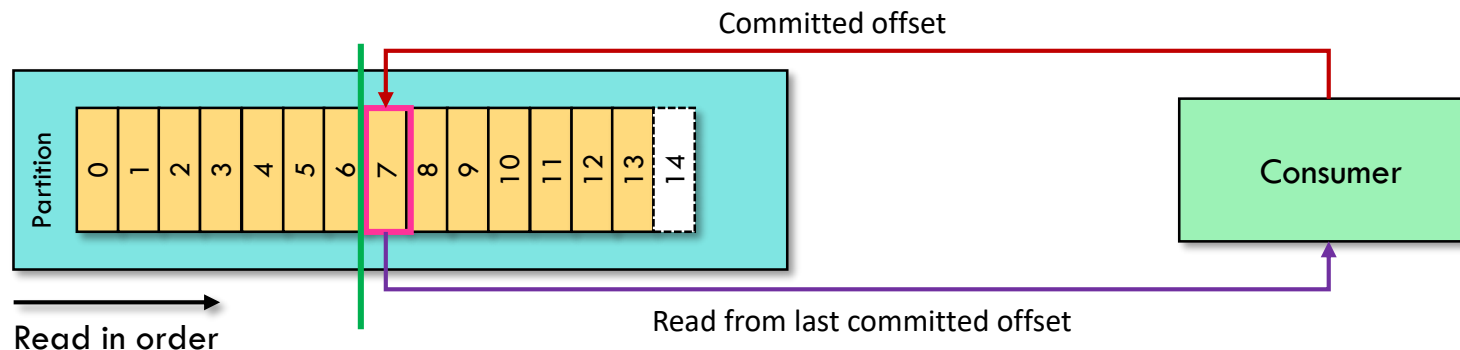
# Consumer Groups

- Consumers read data in **consumer groups**
- Each consumer within a group reads from exclusive partitions
- You cannot have more consumers than partitions (otherwise some will be inactive)



# Consumer Offsets

- Kafka stores the offsets at which a **consumer group** has been reading
- The **offsets** commit live in a Kafka topic named “**\_\_consumer\_offsets**”
- When a consumer has processed data received some Kafka, it should be **committing** the **offsets**
- If a consumer process dies, it will be able to read back from where it left off thanks to consumer offsets!




# Configuring Consumers

- **group.id**

- Kafka consumers are typically part of a **consumer group**. When multiple consumers are subscribed to a topic and belong to the same consumer group, each consumer in the group will receive messages from a different subset of the partitions in the topic.

```
Properties props = new Properties();  
props.put("bootstrap.servers", "broker1:9092,broker2:9092");  
props.put("group.id", "CountryCounter");  
props.put("key.deserializer",  
    "org.apache.kafka.common.serialization.StringDeserializer");  
props.put("value.deserializer",  
    "org.apache.kafka.common.serialization.StringDeserializer");  
  
KafkaConsumer<String, String> consumer = new KafkaConsumer<String,  
String>(props);
```



# Consumers and Consumer Groups

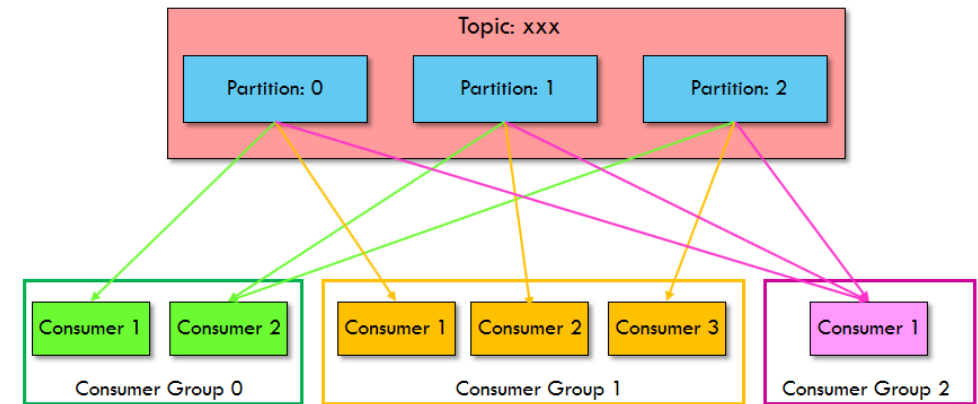
## Create a topic with 4 partitions



```
$ kafka-topics
--create
--zookeeper localhost:2181
--replication-factor 1 --partitions 4
--topic ak03.fourpartition
```

```
root@kafka:/# kafka-topics --create --zookeeper zookeeper:2181 \
> --replication-factor 1 \
> -partitions 4 \
> --topic ak03.fourpartition
WARNING: Due to limitations in metric names, topics with a period ('.') or under
score ('_') could collide. To avoid issues it is best to use either, but not bot
h.
Created topic "ak03.fourpartition"
```

Topic is created!

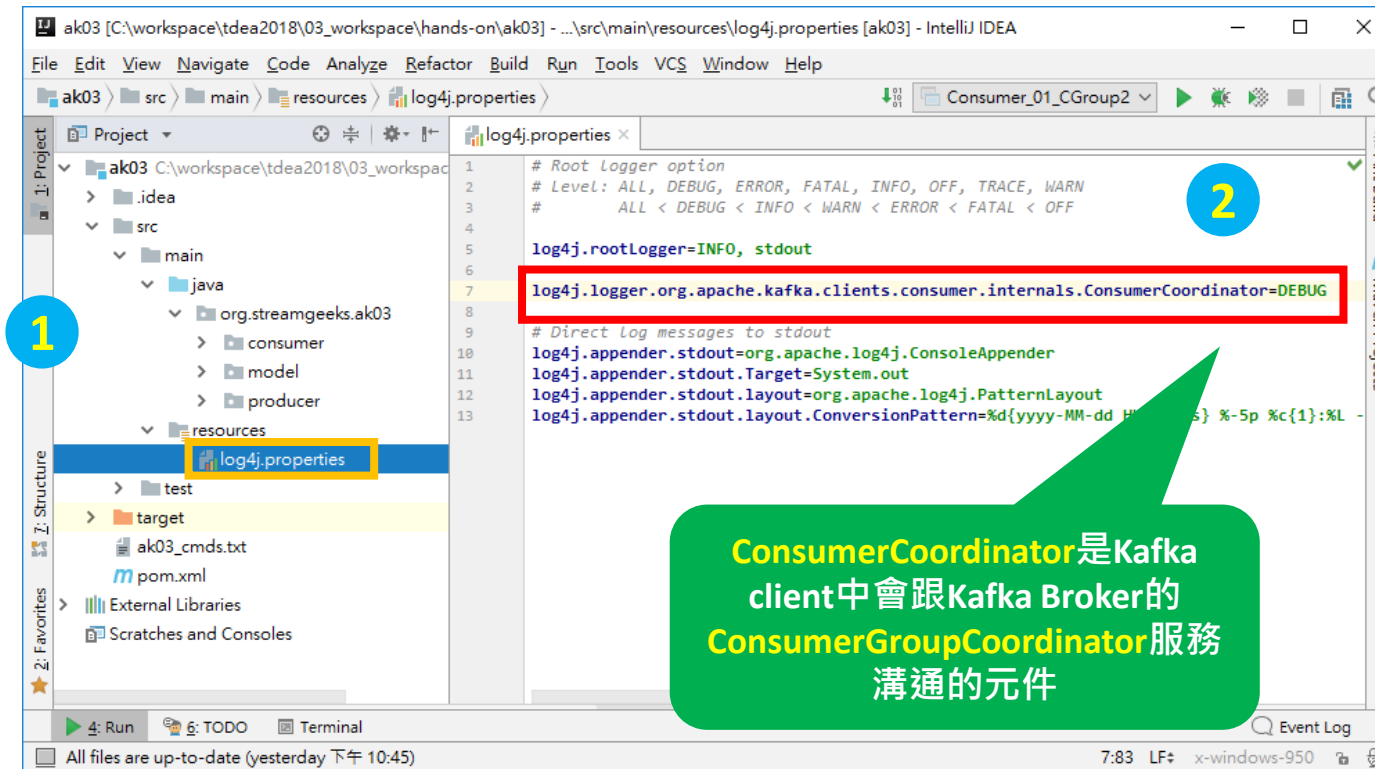


在container裡頭的  
zookeeper服務的  
Hostname



# Consumers and Consumer Groups

## Change ConsumerCoordinator Log Level



# Consumers and Consumer Groups

## Producer\_05\_CGroup : Publish Event:



```
/**
 * 示範：用來持續產生event來給Consumer_01_CGroup進行Rebalance的演示
 */
public class Producer_05_CGroup {
    public static void main(String[] args) {
        // 步驟1. 設定要連線到Kafka集群的相關設定
        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092"); // Kafka集群在那裡?
        props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer"); // 指定msgKey的序列化器
        props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer"); // 指定msgValue的序列化器
        // 步驟2. 產生一個Kafka的Producer的實例
        Producer<String, String> producer = new KafkaProducer<>(props);
        // 步驟3. 指定想要發佈訊息的topic名稱
        String topicName = "ak03.fourpartition";
        int msgCount = 10000; // 1萬筆
        try {
            System.out.println("Start sending messages ...");
            // 步驟4. 產生要發佈到Kafka的訊息 (把訊息封裝進一個ProducerRecord的實例中)
            for(int i=0; i<msgCount; i++) {
                producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i));
                Thread.sleep( millis: 3000); // 註執行緒停個3秒
            }
        } catch (Exception e) {
            // 錯誤處理
            e.printStackTrace();
        }
        // 步驟5. 關掉Producer實例的連線
        producer.close();
        System.out.println("Message sending completed!");
    }
}
```

這個Producer的目的是持續發佈有序列號的訊息來觀察ConsumerGroup的概念!

每3秒發佈一筆訊息!



# Consumers and Consumer Groups

- Producer\_05\_CGroup

```
for(int i=0; i<msgCount; i++) {  
    producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i));  
  
    Thread.sleep( millis: 3000); // 讓主執行緒停個3秒  
}
```

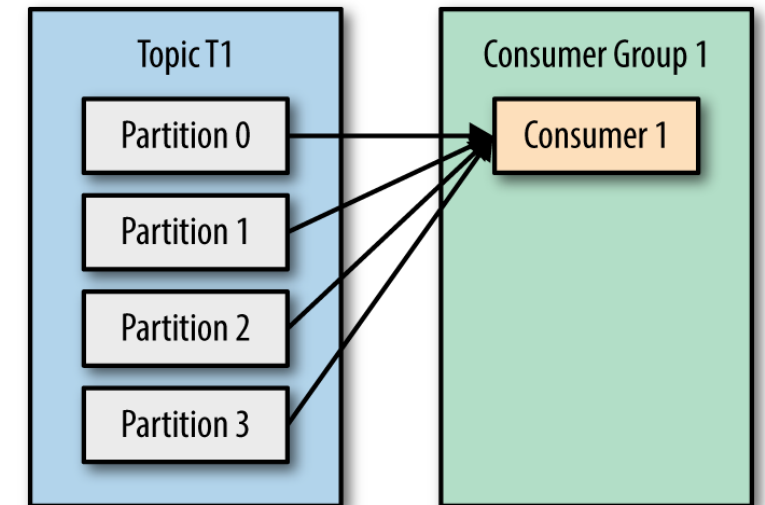
- Consumer\_01\_Cgroup

- One Consumer instance

跑第1個instance, 並  
觀察DEBUG的log

```
INFO ConsumerCoordinator:411 [Consumer clientId=consumer-1, groupId=CG1]  
Revoking previously assigned partitions []  
Setting newly assigned partitions [ak03.fourpartition-0, ak03.fourpartition-2, ak03.fourpartition-1, ak03.fourpartition-3]
```

One Consumer with four partitions





# Consumers and Consumer Groups

- Producer\_05\_CGroup

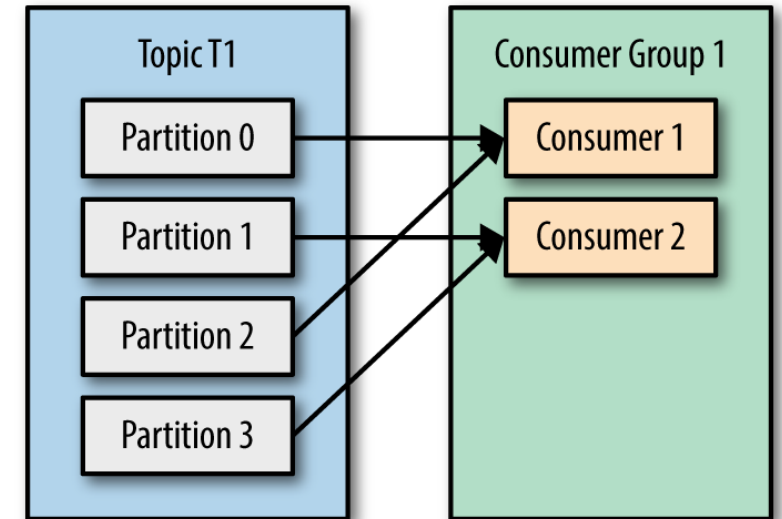
```
for(int i=0; i<msgCount; i++) {  
    producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i));  
  
    Thread.sleep( millis: 3000); // 讓主執行緒停個3秒  
}
```

- Consumer\_01\_Cgroup

- **Two** Consumer instances

跑第2個instance, 並  
觀察DEBUG的log

## Two Consumer with four partitions



```
[Consumer clientId=consumer-1, groupId=CG1] Revoking previously assigned partitions []  
[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-0, ak03.fourpartition-2, ak03.fourpartition-1, ak03.fourpartition-3]  
  
[Consumer clientId=consumer-1, groupId=CG1] Revoking previously assigned partitions [ak03.fourpartition-0, ak03.fourpartition-2, ak03.fourpartition-1, ak03.fourpartition-3]  
[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-0, ak03.fourpartition-1]
```

```
[Consumer clientId=consumer-1, groupId=CG1] Revoking previously assigned partitions []  
[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-2, ak03.fourpartition-3]
```



# Consumers and Consumer Groups

- Producer\_05\_CGroup

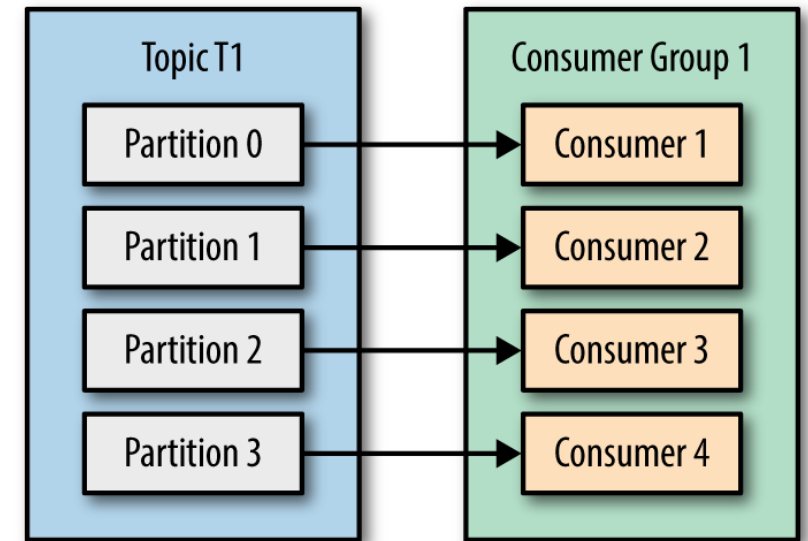
```
for(int i=0; i<msgCount; i++) {  
    producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i));  
  
    Thread.sleep( millis: 3000); // 讓主執行緒停個3秒  
}
```

- Consumer\_01\_Cgroup

- **Four** Consumer instances

跑第4個instance, 並  
觀察DEBUG的log

Three Consumer with four partitions



[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-3]

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-1]

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-0]

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-2]



# Consumers and Consumer Groups

- Producer\_05\_CGroup

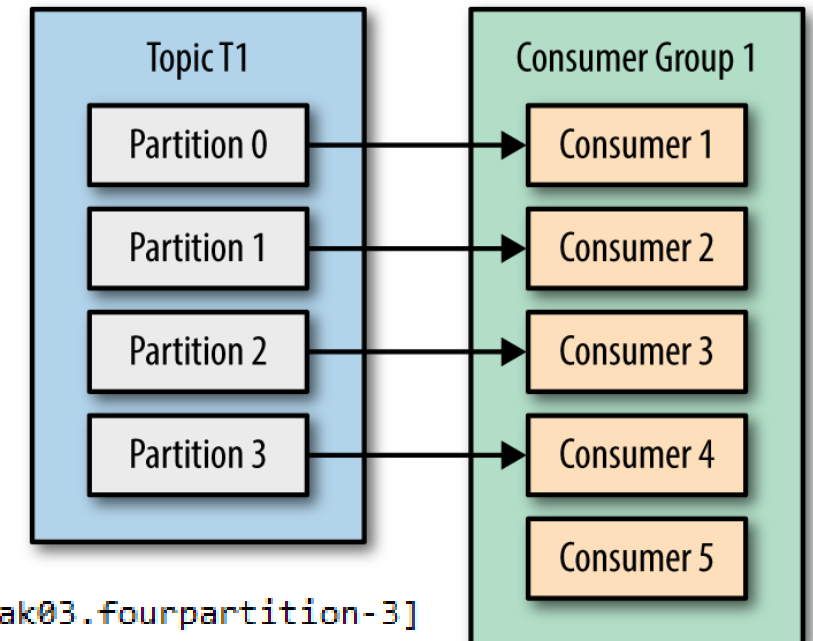
```
for(int i=0; i<msgCount; i++) {  
    producer.send(new ProducerRecord<>(topicName, key: ""+i, value: "msg_"+i));  
  
    Thread.sleep( millis: 3000); // 讓主執行緒停個3秒  
}
```

- Consumer\_01\_Cgroup

- **Five** Consumer instances

跑第**5**個instance, 並  
觀察**DEBUG**的log

## Five Consumer with four partitions



[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-3]

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-1]

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-0]

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions []

[Consumer clientId=consumer-1, groupId=CG1] Setting newly assigned partitions [ak03.fourpartition-2]



# What is Rebalance?

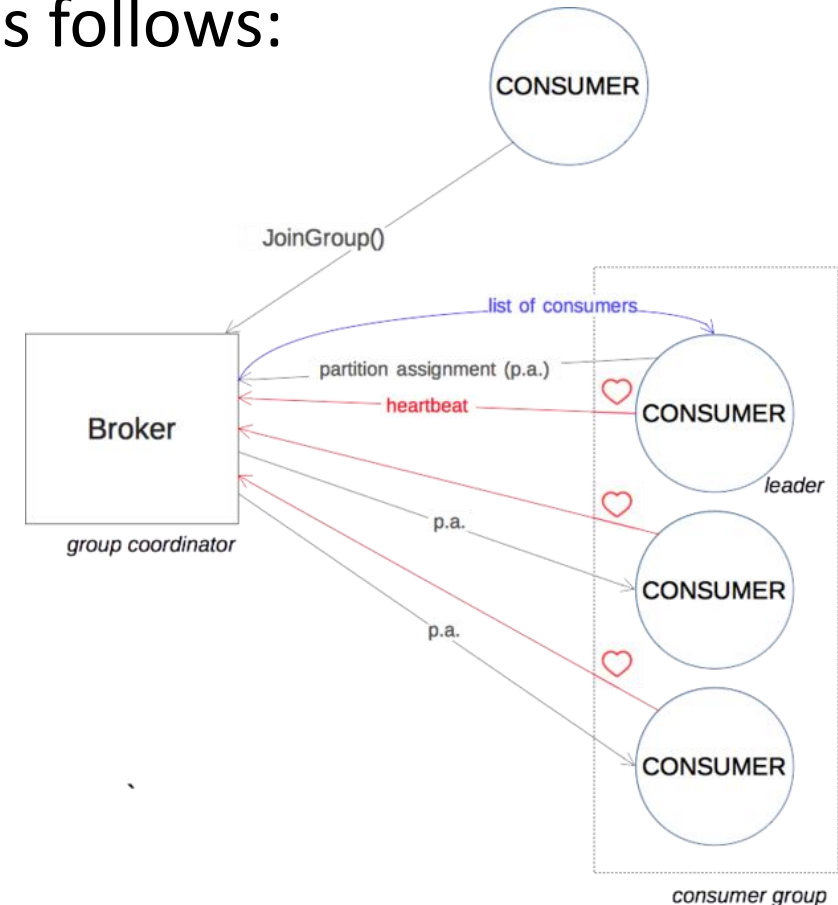
- Every consumer in a **consumer group** is assigned one or more topic partitions exclusively, and **Rebalance** is the re-assignment of partition ownership among consumers.
- A Rebalance happens when:
  - a consumer JOINS the group
  - a consumer SHUTS DOWN cleanly
  - a consumer is considered DEAD by the group coordinator. This may happen after a crash or when the consumer is busy with a long-running processing
  - new partitions are added

# Rebalance process



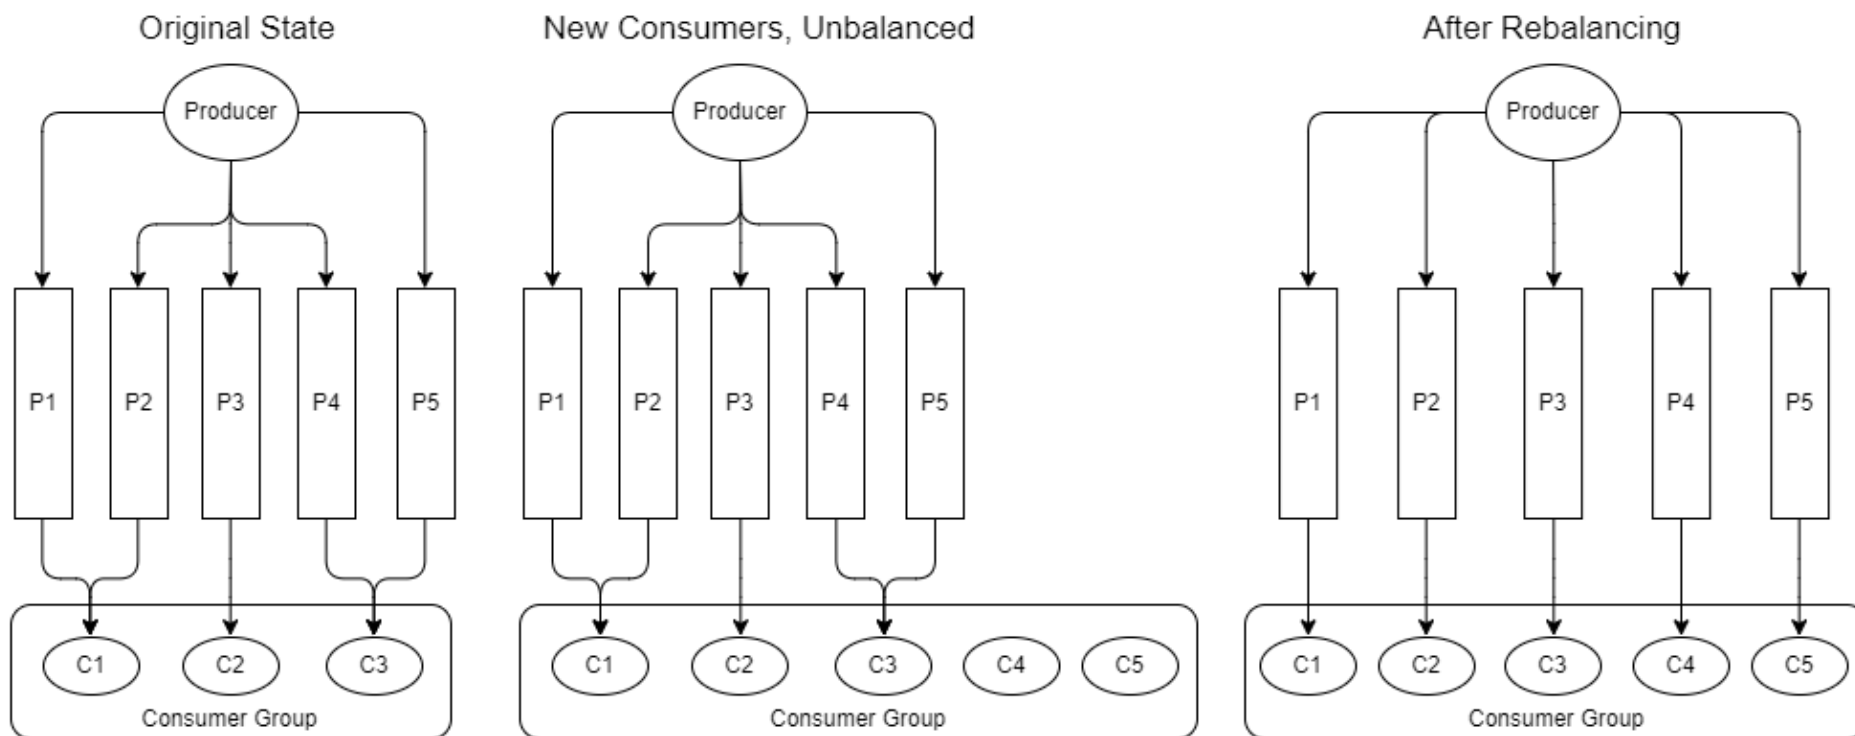
**Rebalance** can be more or less described as follows:

- The leader receives a list of all consumers in the group from the group coordinator and is responsible for assigning a subset of partitions to each consumer.
- After deciding on the partition assignment the group leader sends the list of assignments to the group coordinator, which sends this information to all the consumers.






# Rebalance process



# Commits and Offsets



- Whenever we call **poll()**, broker returns records that consumers in our group have not read yet.
- This means that kafka client tracking which records were read by a consumer of the group.
- We call the action of updating the current position in the partition a **commit**.

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records)  
    {  
        System.out.printf("topic = %s, partition = %s, offset =  
            %d, customer = %s, country = %s\n",  
            record.topic(), record.partition(),
```

# Commits and Offsets

- How does a consumer commit an offset?
  - Kafka client produces a message to Kafka, to a special **\_\_consumer\_offsets** topic, with the committed offset for each partition.
  - If a consumer crashes or a new consumer joins the consumer group, it will trigger a **partition rebalance**.
  - After a **rebalance**, each consumer may be assigned a new set of partitions than the one it processed before.
  - The consumer will read the latest committed offset of each partition and continue from there.

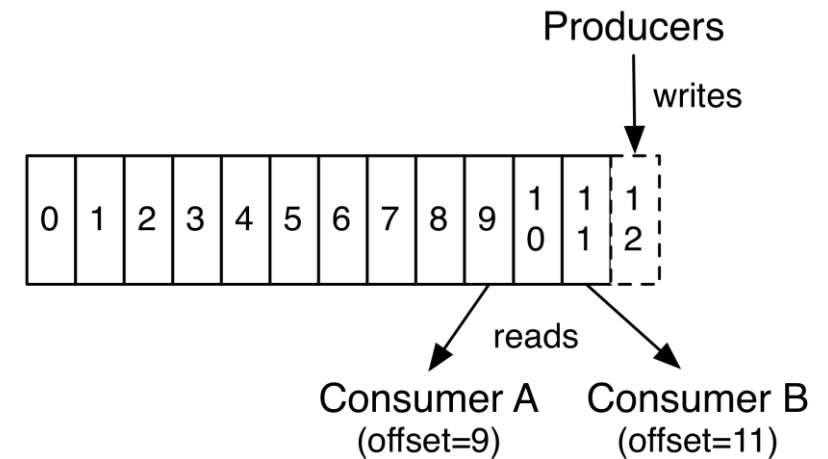
# Commits and Offsets

- Managing offsets has a big impact on the client application.
- The Kafka Consumer API provides multiple ways of committing offsets:
  - Automatic Commit
  - Synchronous Commit (previous **poll()** offset)
  - Asynchronous Commit (previous **poll()** offset)
  - Combining Synchronous and Asynchronous Commits (previous **poll()** offset)
  - Commit Specified Offset

# Configuring Consumers

- **auto.offset.reset**

- This property controls the behavior of the consumer when it starts reading a partition for which it doesn't have a committed offset.
- The default is “**latest**,” which means that lacking a valid offset, the consumer will start reading from the newest records
- The alternative is “**earliest**,” which means that lacking a valid offset, the consumer will read all the data in the partition, starting from the very beginning.



# Commits and Offsets: Automatic Commit

- Configure below configurations:
  - `enable.auto.commit=true`
  - `auto.commit.interval.ms=5000` (default)
- Every **five** seconds the consumer will commit the largest offset your client received from `poll()`.
- Whenever you poll, the consumer checks if it is time to commit, and if it is, it will commit the offsets it returned in the last poll.

# Commits and Offsets: Automatic Commit

- Consumer\_02\_AutoCommit

```
props.put("auto.offset.reset", "earliest"); // 是否從這個ConsumerGroup尚未讀取的partition/offset開始讀  
  
props.put("enable.auto.commit", "true");  
props.put("auto.commit.interval.ms", "5000");
```

▼ resources

log4j.properties

```
log4j.logger.org.apache.kafka.clients.consumer.internals.ConsumerCoordinator=DEBUG
```

```
Sending asynchronous auto-commit of offsets {ak03.fourpartition-0=OffsetAndMetadata{offset=85, metadata=''},  
Committed offset 85 for partition ak03.fourpartition-0  
Committed offset 65 for partition ak03.fourpartition-2  
Committed offset 63 for partition ak03.fourpartition-1  
Committed offset 73 for partition ak03.fourpartition-3  
Completed asynchronous auto-commit of offsets {ak03.fourpartition-0=OffsetAndMetadata{offset=85, metadata=''},
```

# Commits and Offsets: Automatic Commit

- With **auto-commit** enabled, a call to `poll()` *will always commit the last offset returned by the previous poll*. It doesn't know which events were actually processed, so it is critical to always process all the events returned by `poll()` before calling `poll()` again.
- Automatic commits are convenient, but they don't give developers enough control to avoid duplicate messages.



# Commits and Offsets: Manual Commit

- The consumer API has the option of committing the current offset at a point that makes sense to the application developer rather than based on a timer.
- Configure below configurations:
  - **auto.commit.offset=false**
- Offsets will only be committed when the application explicitly chooses to do so.
- The simplest and most reliable of the commit APIs is **commitSync()**.

# Commits and Offsets: Sync Commit

- Consumer\_03\_CommitSync

```
props.put("auto.offset.reset", "earliest"); // 是否從這個ConsumerGroup尚未讀取的partition/offset開始讀  
props.put("enable.auto.commit", "false");
```

▼ resources

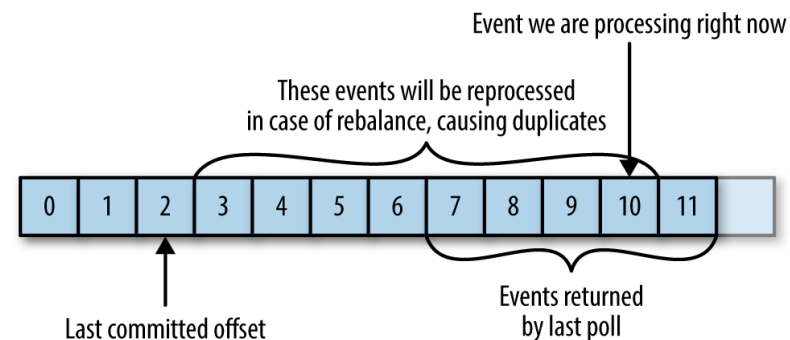
log4j.properties

log4j.logger.org.apache.kafka.clients.consumer.internals.ConsumerCoordinator=DEBUG

```
while (true) {  
    // 請求Kafka把新的訊息吐出來  
    ConsumerRecords<String, String> records = consumer.poll( timeout: 1000);  
  
    // 如果有任何新的訊息就會進到下面的迭代  
    for (ConsumerRecord<String, String> record : records){  
        // ** 在這裡進行商業邏輯與訊息處理 **  
        // ....  
    }  
    try {  
        consumer.commitSync(); // <-- Commit Last Poll()的offset  
    } catch (CommitFailedException e) {  
        // 錯誤處理  
        System.out.println("commit failed");  
        e.printStackTrace();  
    }  
}
```

# Commits and Offsets: Sync Commit

- This **commitSync()** will commit the latest offset returned by **poll()** and return once the offset is committed, throwing an exception if commit fails for some reason.
- Make sure you call **commitSync()** after you are done processing all the records in the collection.
- When **rebalance** is triggered, all the messages from the beginning of the most recent batch until the time of the rebalance will be processed twice.



# Commits and Offsets: Asynchronous Commit

- One drawback of manual commit is that the application is **blocked** until the broker responds to the commit request.
- Another option is the **asynchronous** commit API. Instead of waiting for the broker to respond to a commit, we just send the request and continue on:

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records)  
    {  
        System.out.printf("topic = %s, partition = %s,  
offset = %d, customer = %s, country = %s\n",  
record.topic(), record.partition(), record.offset(),  
record.key(), record.value());  
    }  
    consumer.commitAsync(); ❶  
}
```

# Commits and Offsets: Asynchronous Commit

- Consumer\_04\_CommitAsync

```
props.put("auto.offset.reset", "earliest"); // 是否從這個ConsumerGroup尚未讀取的partition/offset開始讀  
props.put("enable.auto.commit", "false");
```

▼ resources

log4j.properties

log4j.logger.org.apache.kafka.clients.consumer.internals.ConsumerCoordinator=DEBUG

```
while (true) {  
    // 請求Kafka把新的訊息吐出來  
    ConsumerRecords<String, String> records = consumer.poll( timeout: 1000);  
    // 如果有任何新的訊息就會進到下面的迭代  
    for (ConsumerRecord<String, String> record : records){  
        // ** 在這裡進行商業邏輯與訊息處理 **  
        // ..  
    }  
    consumer.commitAsync(new OffsetCommitCallback() {  
        @Override  
        public void onComplete(Map<TopicPartition, OffsetAndMetadata> offsets, Exception e) {  
            if (e != null) {  
                // 錯誤處理  
                System.out.println("commit failed");  
                e.printStackTrace();  
            }  
        }  
    });  
};
```

# Commits and Offsets: Asynchronous Commit

- The drawback is that while **commitSync()** will retry the commit until it either succeeds or encounters a non-retriable failure
- **commitAsync()** will not retry.
- **commitAsync()** gives you an option to pass in a callback that will be triggered when the broker responds.

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records) {  
        System.out.printf("topic = %s, partition = %s,  
offset = %d, customer = %s, country = %s\n",  
record.topic(), record.partition(), record.offset(),  
record.key(), record.value());  
    }  
    consumer.commitAsync(new OffsetCommitCallback() {  
        public void onComplete(Map<TopicPartition,  
OffsetAndMetadata> offsets, Exception exception) {  
            if (e != null)  
                log.error("Commit failed for offsets {}", offsets, e);  
        }  
    }); ❶
```

# Commits and Offsets: Combining Sync and Async Commits

```
try {
    while (true) {
        ConsumerRecords<String, String> records = consumer.poll(100);
        for (ConsumerRecord<String, String> record : records) {
            System.out.printf("topic = %s, partition = %s, offset = %d,
                               customer = %s, country = %s\n",
                               record.topic(), record.partition(),
                               record.offset(), record.key(), record.value());
        }
        consumer.commitAsync(); ❶
    }
} catch (Exception e) {
    log.error("Unexpected error", e);
} finally {
    try {
        consumer.commitSync(); ❷
    } finally {
        consumer.close();
    }
}
```

# Commits and Offsets: Combining Sync and Async Commits

- Consumer\_05\_CommitSyncAsync

```
// 步驟5. 持續的拉取Kafka有進來的訊息
try {
    System.out.println("Start listen incoming messages ...");

    while (true) {
        // 請求Kafka把新的訊息吐出來
        ConsumerRecords<String, String> records = consumer.poll( timeout: 1000);

        // 如果有任何新的訊息就會進到下面的迭代
        for (ConsumerRecord<String, String> record : records) {
            // ** 在這裡進行商業邏輯與訊息處理 **
            // ...
        }
        consumer.commitAsync();
    }
} catch (Exception e) {
    // log.error("Unexpected error", e);
    e.printStackTrace();
} finally {
    // 步驟6. 如果收到結束程式的訊號時關掉Consumer實例的連線
    try {
        consumer.commitSync();
    } finally {
        consumer.close();
    }
}
System.out.println("Stop listen incoming messages");
}
```



# Commits and Offsets: Commit Specified Offset

- Committing the latest offset only allows you to commit as often as you finish processing batches.
- What if **poll()** returns a huge batch and you want to commit offsets in the middle of the batch to avoid having to process all those rows again if a rebalance occurs?
- You can't just call **commitSync()** or **commitAsync()** —this will commit the last offset returned, which you didn't get to process yet.
- Consumer API allows you to call **commitSync()** and **commitAsync()** and pass a map of **partitions** and **offsets** that you wish to commit.

# Commits and Offsets: Commit Specified Offset

```
private Map<TopicPartition, OffsetAndMetadata> currentOffsets =  
    new HashMap<>(); ❶  
int count = 0;  
  
....  
  
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records)  
    {  
        System.out.printf("topic = %s, partition = %s, offset = %d,  
            customer = %s, country = %s\n",  
            record.topic(), record.partition(), record.offset(),  
            record.key(), record.value()); ❷  
        currentOffsets.put(new TopicPartition(record.topic(),  
            record.partition()), new  
            OffsetAndMetadata(record.offset()+1, "no metadata")); ❸  
        if (count % 1000 == 0) ❹  
            consumer.commitAsync(currentOffsets, null); ❺  
        count++;  
    }  
}
```

# Commits and Offsets: Commit Specified Offset

- Consumer\_06\_CommitSpecified

```
// 步驟5. 產生一個紀錄TopicPartition與Offset的容器
Map<TopicPartition, OffsetAndMetadata> currentOffsets = new HashMap<>();
// 步驟6. 持續的拉取Kafka有進來的訊息
int recordCount = 0;
try {
    System.out.println("Start listen incoming messages ...");
    while (true) {
        // 請求Kafka把新的訊息吐出來
        ConsumerRecords<String, String> records = consumer.poll( timeout: 1000);
        // 如果有任何新的訊息就會進到下面的迭代
        for (ConsumerRecord<String, String> record : records) {
            // ** 在這裡進行商業邏輯與訊息處理 **
            // 取出相關的metadata
            String topic = record.topic();
            int partition = record.partition();
            long offset = record.offset();
            TimestampType timestampType = record.timestampType();
            long timestamp = record.timestamp();
            // 取出msgKey與msgValue
            String msgKey = record.key();
            String msgValue = record.value();
            // 秀出metadata與msgKey & msgValue訊息
            System.out.println(topic + "-" + partition + "-" + offset + " : (" + record.key() + ", " + record.value() + ")");
            // 紀錄現在的TopicPartition與Offset
            currentOffsets.put(new TopicPartition(topic, partition), new OffsetAndMetadata( offset: offset+1, metadata: "no metadata"));
            // 決定何時要commit
            if (recordCount % 1000 == 0) {
                Map<TopicPartition, OffsetAndMetadata> offsetToCommits = copyHashmap(currentOffsets);
                consumer.commitAsync(offsetToCommits, callback: null); // 對特定的offset進行commit
                currentOffsets.clear(); // 清除原有TopicPartition與Offset的紀錄
            }
            recordCount++;
        }
    }
}
consumer.commitAsync();
```

# Rebalance Listeners

- If you know your consumer is about to lose ownership of a partition, you will want to commit offsets of the last event you've processed.
- The consumer API allows you to run your own code when partitions are added or removed from the consumer.
- You do this by passing a **ConsumerRebalanceListener** when calling the subscribe() method
- **ConsumerRebalanceListener** has two methods you can implement:
  - onPartitionsRevoked
  - onPartitionsAssigned

# Commits and Offsets: RebalanceListener

- Consumer\_07\_CommitSpecified

```
public class RebalanceHandler implements ConsumerRebalanceListener {  
    private Map<TopicPartition, OffsetAndMetadata> currentOffsets = new HashMap<>();  
    private Consumer consumer;  
  
    public RebalanceHandler(Consumer consumer, Map<TopicPartition, OffsetAndMetadata> currentOffsets){  
        this.consumer = consumer;  
        this.currentOffsets = currentOffsets;  
    }  
  
    // 當Rebalance被觸發後, Consumer被取回被assigned的partitions  
    @Override  
    public void onPartitionsRevoked(Collection<TopicPartition> partitions) {  
        System.out.println("Lost partitions in rebalance. Committing current Offsets: " + currentOffsets);  
        consumer.commitSync(currentOffsets);  
    }  
  
    // 當Rebalance被觸發後, Consumer被通知有那些partition被assigned  
    @Override  
    public void onPartitionsAssigned(Collection<TopicPartition> partitions) {  
  
    }  
}
```

# Reference

1. Kafka: The Definitive Guide – O'REILLY



2. Kafka In Action - MANNING



3. Learn Apache Kafka for Beginners – Udemy (Stephane Maarek)



4. Confluent Document of Kafka – confluent.io



